

Palaeodata wrangling with the Tidyverse



Steve Juggins



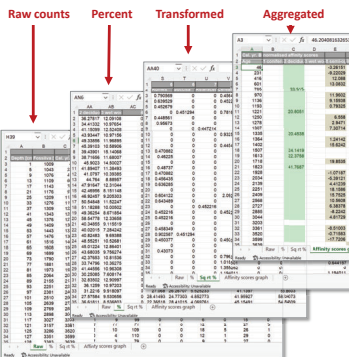
Guiding principles

- **Everything** is possible in **Base R** but solutions may be ad-hoc and need a mix functions & packages with different syntax: the Wild West that is Base R!
- **Tidyverse** offers a consistent approach to data wrangling, fast, concise & logical syntax, and very flexible.
- Some things may be **faster** in Base R but speed is rarely a limiting factor.
- **Focus** on code that is simple, readable and easy to understand (promote learning, easier debugging, and reproducibility).
- **Today** - don't get bogged down in the details of individual functions and operations. Focus on the big picture - the details will come with practice.

3

Do you need the tidyverse?

Do you have many separate files / sheets of the same data? (counts, percentages, subsets, aggregated)



If yes, then you should aim to have only one file and do all processing in R.

- More **efficient** (fewer files to keep track of)
- Reduces **errors** (and easier to correct errors and update data – only 1 file to change)
- **Reproducible** (record of what you did and why)

Day 1 aims

Morning

Aims: To reproduce Woodbridge et al. Figure 3 using Tidyverse methods & explore changes in palynological richness.

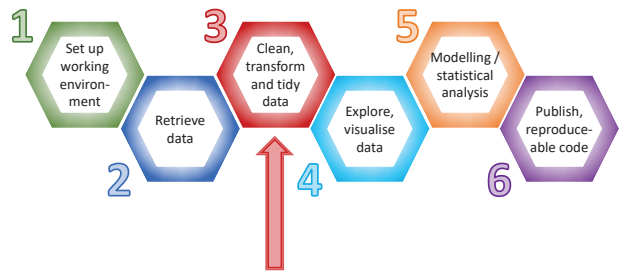
1. Work through analysis of a single site using base R
2. Repeat using Tidyverse
3. Apply to all (41) sites & create synthesis figure
4. Modelling species richness with the Tidyverse

Afternoon

- Demonstration of my new package **riojaPlot** for plotting stratigraphic diagrams
- Work with your own data, ask questions, get help

7

Palaeodata Workflow



Not sexy, not clever, not always straightforward, and can be very time consuming.
Having a good set of tools in R can:

- Increase efficiency
- Reduce errors
- Self-documenting & reproducible

Based on Fig 3, Slater et al. 2010. Using R in hydrology: a review of recent developments and future directions. Hydrology and Earth System Sciences 23:2939–2963.

2

Resources

Whickham & Grolemond, R for Data Science, free online <https://r4ds.had.co.nz/>

Locke, R Fundamentals 2: Data Manipulation (£11) <https://itsalocke.com/company/books/>

Bruno Rodrigues, Modern R with the tidyverse free online https://b-rodrigues.github.io/modern_R/ or buy pdf for < £10

Murray Logan's website, excellent tutorials & book <https://www.flutterbys.com.au/stats/>


Rstudio cheatsheets <https://www.rstudio.com/resources/cheatsheets/>
stackoverflow + 1001 other websites...

4

Example: Impact of the Neolithic Agricultural Transition

Jessie Woodbridge, Ralph Fyfe & colleagues
Woodbridge et al. (2014) & Fyfe et al. (2010)

Aim: Reconstruct anthropogenic land-cover change in Britain over the last 9000 years



Land-cover reconstructions for Britain
Figure 3 from Woodbridge et al. (2014)




Thanks to Jessie Woodbridge and Ralph Fyfe for sharing their data!

6

Example: Impact of the Neolithic Agricultural Transition


1. Extract dated pollen profiles from European pollen database (+ other sources), 41 sites in total.




2. Classify each pollen type to a land cover class (LCC)



3. For each site, transform species data to major vegetation types (land cover class: LCC) and calculate the dominant LCC for each sample.



4. Sum dominant LCCs for all sites in 200 year time slices, convert to % of each LCC & plot.



The data

1. **Woodbridge_et_al_2014_Data.xlsx**: Excel file with pollen counts for 42 sites, each site on a separate tab. Columns for pollen types, Depth and Cal. Age BP + alternative chronologies, sample IDs & and pollen sum.
2. **LCC_Info.xlsx**: Excel file with the following worksheets:

LCC_lookup: Lookup table of taxon names and corresponding land cover classes.


Site_info: Information about each site, including region, grid references etc.

9

Land Cover Classification (LCC) method

Step1: Calculate LCC for each site

1. Import data
2. Clean data - remove unwanted columns
3. Rename columns
4. Convert counts to percentages
5. Allocate each pollen type to an LCC class and sum sqrt-percentages across each class
6. Normalise the sqrt-percent LCC data to 100%
7. Identify the dominant LCC class at each level



Step2: Aggregate multiple sites

8. Aggregate data from all sites and count the number of levels in each LCC class in 200 year time slices (N)
9. Convert N to percentage of each LCC class in each time slice
10. Plot the aggregated data


10

Part 1: Code Step 1 in Base R

Palaeodata wrangling with the tidyverse

What is the Tidyverse?

- Collection of R packages designed for data science that share common interface standards, grammar and data structures
- You spend more time on concepts and less sorting out syntax
- Promotes code readability



- Both **general purpose** packages: **tibble** for data frames, **tidyr** for tidying, **dplyr** for transforming and **ggplot2** for visualisation, and **specialised** packages for data import (**readr**), dates & time (**lubridate**, **hms**), strings (**stringr**), factors (**forcats**) plus others.

13

Core features

1. Uses **tibbles** rather than **data frames**
refined **print** method
strict about **\$** subsetting, doesn't like **rownames**
easier to create **nested** data frames
2. Encourages use of the pipe **%>%**
Instead of **summary(lm(y~x, data=df))**
Use **df %>% lm(y~x, data=.) %>% summary()**
3. Promotes use of **tidy** data
Every column is a variable
Every row is an observation
Every cell is a single value

14

Tidy data

Not tidy !

(because the variable is pollen type and the attributes are taxon and count)

Depth (cm)	Cal. yr. BP	Pinus sylv	Taxus bacc	Betula	Betula/Cor
132	1514	3	1	40	2
164	1618.5	3	3	33	1
196	1723	3	2	34	2
228	1827.5	4	4	33	1
260	1932	5	3	34	3
268	2143.75	4	3	43	0
276	2355.5	14	2	40	0
284	2567.25	11	1	60	0
292	2779	10	3	54	2

Wide format

Convert from **wide** to **long** with **pivot_longer()**

```
df %>% pivot_longer(
  cols=c("Depth (cm)", "Cal. yr. BP"),
  names_to="Taxon", values_to="Count")
```

Tidy !

(Each column is a variable, pollen counts are described by key/value pairs)

Depth (cm)	Cal. yr. BP	Taxon	Count
132	1514	Pinus sylvestris	3
132	1514	Taxus baccata	1
132	1514	Betula	40
132	1514	Betula/Corylus/Myrica	2
164	1618.5	Pinus sylvestris	3
164	1618.5	Taxus baccata	3
164	1618.5	Betula	33
164	1618.5	Betula/Corylus/Myrica	1
196	1723	Pinus sylvestris	3

Long format

Convert from **long** to **wide** with **pivot_wider()**

```
df %>% pivot_wider(
  id_cols=c("Depth (cm)", "Cal. yr. BP"),
  names_from="Taxon", values_from="Count")
```

In package **tidyr**

15

dplyr functions for data transformation

- select**: select columns (**subset**) dplyr function
Base R equivalent
- filter**: subsetting rows based on condition (**subset**)
- arrange**: sorting (**sort**, **order**)
- distinct**: find unique values (**unique**)
- slice**: selecting rows based on position (**[]**)
- mutate**: create new columns (**transform**)
- group_by**: defining groups of rows to process subsets
- summarise**: summarise data (optionally by group) (**aggregate**)
- *_join**: merging data sets (**merge**)
- bind_rows**, **bind_cols**: combine multiple dfs by row or column (**rbind**, **cbind**)

16

Rewriting code as a function

- To apply our LCC function to multiple sites we can use a loop and collate results at each iteration.
- We could include all the LCC code within the loop but this is not good – makes the code more **complex**, and we should aim to **reuse** code, not **repeat** it.
- Encapsulate the code in a **function**, and call the function from within the loop.

Defining a function:

```
my_fun <- function(arg1, arg2, ...) {  
  statements  
  return(object)  
}
```

Example:

```
my_fun <- function(a, b) {  
  c <- a + b  
  return(c)  
}  
my_fun(2, 3)  
[1] 5
```


25

Working with multiple datasets – Tidyverse

Two approaches with Tidy functions that avoid loops and produce more readable, elegant code:

1. Convert each df to long format, and stack, and perform summaries / transformations by grouping of Site and Age/Depth
2. Convert to nested data frame that stores data of each site within the cell of a larger organising table. Apply a function to each nested table and collate results.

We use second approach to clean data and convert to long format and stack the data, then use first approach to calculate LCC from the stacked dataset.



27

Part 3: Code to aggregate multiple sites Base R

Working with nested data using purrr

Useful functions:

tidyr::nest(df, cols): creates data frame with cols nested within grouped defined by the non-nesting columns.

tidyr::unnest(df, cols): unnests a nested column col of a nested data frame.

purrr::map applies a function to each element of a list or vector (ie. each data frame of a nested data frame)

28

Working with nested data frames

Original data frame

```
iris  
A tibble: 150 x 5  
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species  
    <dbl>         <dbl>         <dbl>         <dbl> <fct>  
1     5.1         3.5         1.4         0.2 setosa  
2     4.9         3         1.5         0.2 setosa  
3     4.7         3.2         1.3         0.2 setosa  
4     4.6         3.1         1.5         0.2 setosa  
5     5           3.6         1.4         0.2 setosa  
6     5.4         3.9         1.7         0.2 setosa  
7     4.6         3.4         1.4         0.3 setosa  
8     5           3.4         1.5         0.2 setosa  
9     4.4         2.9         1.4         0.2 setosa  
10    4.9         3.1         1.5         0.1 setosa
```

Nested data frame

```
n_iris <- iris %>% nest(data = ~Species)  
n_iris  
A tibble: 3 x 2  
  Species data  
    <fct>   <list>  
1 setosa  <tibble [50 x 4]>  
2 versicolor <tibble [50 x 4]>  
3 virginica <tibble [50 x 4]>
```

List-column named "data"

Use map to apply function summary to each nested table

```
n_iris %>% summarise(summary = map(summary))  
[[1]]  
  Sepal.Length Sepal.Width Petal.Length Petal.Width  
Min.   14.300 Min.   12.300 Min.   11.000 Min.   10.100  
1st Qu.:4.800 1st Qu.:3.200 1st Qu.:1.400 1st Qu.:0.200  
Median :5.000 Median :3.400 Median :1.500 Median :0.200  
Mean   :5.006 Mean   :3.428 Mean   :1.462 Mean   :0.246  
3rd Qu.:5.200 3rd Qu.:3.675 3rd Qu.:1.575 3rd Qu.:0.300  
Max.   :5.800 Max.   :4.400 Max.   :1.900 Max.   :0.600  
[[2]]  
  Sepal.Length Sepal.Width Petal.Length Petal.Width  
Min.   14.900 Min.   12.000 Min.   13.000 Min.   11.000
```

29

Working with nested data frames

Apply base R function **summary** to each nested table and add results to original nested df

```
n_iris %>% mutate(summary = map(summary))  
A tibble: 3 x 3  
  Species data summary  
    <fct>   <list>   <list>  
1 setosa  <tibble [50 x 4]> <tibble [6 x 4]>  
2 versicolor <tibble [50 x 4]> <tibble [6 x 4]>  
3 virginica <tibble [50 x 4]> <tibble [6 x 4]>
```

Use formula version to supply arguments to function (in this case **lm**)

```
n_iris %>% mutate(model = map(summary, ~lm(Sepal.Width~Sepal.Length, data=.x)))  
A tibble: 3 x 3  
  Species data model  
    <fct>   <list>   <list>  
1 setosa  <tibble [50 x 4]> <lm>  
2 versicolor <tibble [50 x 4]> <lm>  
3 virginica <tibble [50 x 4]> <lm>
```


Use ".x" to pass the df in the first argument to map to the mapped function

30


Part 4: Code to aggregate multiple sites using Tidyverse

Part 5: Modelling with the Tidyverse

- Calculate palynological richness using rarefaction



- Apply model to all sites



Very simple questions:
Which sites show significant trend in richness?
How does richness vary with landscape "openness"

31

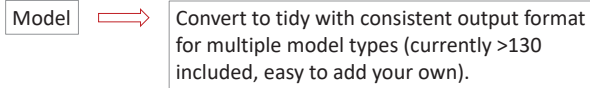
32

broom for tidying model output

The problem



The **broom** solution



- tidy:** tibble that summarises model findings (e.g. coefficients, p-values)
- glance:** concise one-row summary (e.g. r-squared, degf)
- augment:** columns original data was modelled on (e.g. predictions, residuals)

33

Afternoon (13:30-15:30): Bring your own data

